

h1

October 17, 2019

Christian Steinmetz

Generate a dataset of two-dimensional points, and choose a random line in the plane as your target function f , where one side of the line maps to $+1$ and the other side to -1 . Let the inputs $\mathbf{x}_n \in \mathbb{R}^2$ be random points in the plane, and evaluate the target function f on each \mathbf{x}_n to get the corresponding output $y_n = f(\mathbf{x}_n)$. Experiment with the perceptron algorithm in the following settings:

```
[1]: import numpy as np
import matplotlib.pyplot as plt
#import matplotlib as mpl
#mpl.rcParams['figure.dpi'] = 300
%config InlineBackend.figure_format = 'retina'

pColor = 'b'
nColor = 'r'
```

```
[2]: def test_convergence(w, x, y):
    for x_n, y_n in zip(x, y):

        x_n = np.insert(x_n, 0, 1)

        if y_n > 0 and np.dot(w, x_n) < 0:
            return False

        if y_n < 0 and np.dot(w, x_n) >= 0:
            return False

    return True
```

1 a

Generate a dataset of size 20. Plot the examples $\{(\mathbf{x}_n, y_n)\}$ as well as the target function f on a plane.


```
[3]: # generate a dataset of points in 2D space from -1 to 1
N = 20
x = (np.random.rand(N,2) * 2) - 1
```

```
[4]: # create a "random" target function
w_t = [0, 0.5, -0.5]
def target(x_n):
    return 1 if np.dot(w_t, np.insert(x_n, 0, 1)) >= 0 else -1

[5]: # create vector of y values (labels) with target function
y = [target(x_n) for x_n in x]

[6]: plt.scatter(x[:,0], x[:,1], c=[pColor if y_n >= 0 else nColor for y_n in y]) #
    ↪ plot the points in the dataset
plt.plot(np.linspace(-2,2,50), -(w_t[1]/w_t[2]) * (np.linspace(-2,2,50)
    ↪ -(w_t[0]/w_t[2])), c='k')
plt.axis('equal')
plt.axis([-1.25, 1.25, -1.25, 1.25])
```

```
[6]: [-1.25, 1.25, -1.25, 1.25]
```



2 b

Run the perceptron algorithm on the dataset. Report the number of updates that the algorithm takes before converging. Plot the examples $\{(\mathbf{x}_n, y_n)\}$, the target function f , and the final hypothesis g in the same figure.

```
[7]: # randomly initialize weights
w = np.zeros(3)
```

```

idx = 0

while not test_convergence(w, x, y):
    # pick a random sample
    n = np.random.randint(0,x.shape[0]-1)
    x_n = np.insert(x[n,:], 0, 1)
    y_n = y[n]

    if y_n > 0 and np.dot(w, x_n) < 0:
        w += x_n

    if y_n < 0 and np.dot(w, x_n) >= 0:
        w -= x_n

    idx += 1

print(f"Converged in {idx} steps with w = {w}")


```

Converged in 17 steps with w = [0. 1.43335106 -1.13837864]

```

[8]: plt.scatter(x[:,0], x[:,1], c=[pColor if y_n >= 0 else nColor for y_n in y]) #
    ↪ plot the points in the dataset
plt.plot(np.linspace(-2,2,50), -(w_t[1]/w_t[2]) * (np.linspace(-2,2,50)
    ↪ -(w_t[0]/w_t[2])), c='k')
plt.plot(np.linspace(-2,2,50), -(w[1]/w[2]) * (np.linspace(-2,2,50) -(w[0]/
    ↪ w[2])), c='g')
plt.axis('equal')
plt.axis([-1.25, 1.25, -1.25, 1.25])
plt.show()

```



3 c

Repeat everything in b) with another randomly generated dataset of size 20, and compare the result to b).

```
[9]: # generate a dataset of points in 2D space from -1 to 1
N = 20
x = (np.random.rand(N,2) * 2) - 1

# create vector of y values (labels) with target function
y = [target(x_n) for x_n in x]

# randomly initialize weights
w = np.zeros(3)
idx = 0

while not test_convergence(w, x, y):
    # pick a random sample
    n = np.random.randint(0,x.shape[0]-1)
    x_n = np.insert(x[n,:], 0, 1)
    y_n = y[n]

    if y_n > 0 and np.dot(w, x_n) < 0:
        w += x_n
```

```

if y_n < 0 and np.dot(w, x_n) >= 0:
    w -= x_n


idx += 1

print(f"Converged in {idx} steps with w = {w}")

plt.scatter(x[:,0], x[:,1], c=[pColor if y_n >= 0 else nColor for y_n in y]) #_
    ↪plot the points in the dataset
plt.plot(np.linspace(-2,2,50), -(w_t[1]/w_t[2]) * (np.linspace(-2,2,50) -_
    ↪(w_t[0]/w_t[2])), c='k')
plt.plot(np.linspace(-2,2,50), -(w[1]/w[2]) * (np.linspace(-2,2,50) - (w[0]/_
    ↪w[2])), c='g')
plt.axis('equal')
plt.axis([-1.25, 1.25, -1.25, 1.25])
plt.show()

```

Converged in 53 steps with w = [0. 1.22477878 -1.17061912]



4 d

Repeat everything in b) with another randomly generated dataset of size 100, and compare the result to b).

```

[10]: # generate a dataset of points in 2D space from -1 to 1
N = 100
x = (np.random.rand(N,2) * 2) - 1

# create vector of y values (labels) with target function
y = [target(x_n) for x_n in x]

# randomly initialize weights
w = np.zeros(3)
idx = 0

while not test_convergence(w, x, y):
    # pick a random sample
    n = np.random.randint(0,x.shape[0]-1)
    x_n = np.insert(x[n,:], 0, 1)
    y_n = y[n]

    if y_n > 0 and np.dot(w, x_n) < 0:
        w += x_n

    if y_n < 0 and np.dot(w, x_n) >= 0:
        w -= x_n


    idx += 1

print(f"Converged in {idx} steps with w = {w}")

plt.scatter(x[:,0], x[:,1], c=[pColor if y_n >= 0 else nColor for y_n in y]) #
    ↳ plot the points in the dataset
plt.plot(np.linspace(-2,2,50), -(w_t[1]/w_t[2]) * (np.linspace(-2,2,50) -
    ↳ (w_t[0]/w_t[2])), c='k')
plt.plot(np.linspace(-2,2,50), -(w[1]/w[2]) * (np.linspace(-2,2,50) - (w[0]/
    ↳ w[2])), c='g')
plt.axis('equal')
plt.axis([-1.25, 1.25, -1.25, 1.25])
plt.show()

```

Converged in 217 steps with w = [0. 3.47857187 -3.55081271]



5 e

Repeat everything in b) with another randomly generated dataset of size 1000, and compare the result to b).

```
[11]: # generate a dataset of points in 2D space from -1 to 1
N = 1000
x = (np.random.rand(N,2) * 2) - 1

# create vector of y values (labels) with target function
y = [target(x_n) for x_n in x]

# randomly initialize weights
w = np.zeros(3)
idx = 0

while True:
    # pick a random sample
    n = np.random.randint(0,x.shape[0]-1)
    x_n = np.insert(x[n,:], 0, 1)
    y_n = y[n]

    if y_n > 0 and np.dot(w, x_n) < 0:
        w += x_n
```

```

if y_n < 0 and np.dot(w, x_n) >= 0:
    w -= x_n


# check for converge every 100 steps
if idx % 100 == 0:
    if test_convergence(w, x, y):
        break
idx += 1

print(f"Converged in {idx} steps with w = {w/np.linalg.norm(w, ord=1)}")

plt.scatter(x[:,0], x[:,1], c=[pColor if y_n >= 0 else nColor for y_n in y]) #_
    ↳plot the points in the dataset
plt.plot(np.linspace(-2,2,50), -(w_t[1]/w_t[2]) * (np.linspace(-2,2,50) -_
    ↳(w_t[0]/w_t[2])), c='k')
plt.plot(np.linspace(-2,2,50), -(w[1]/w[2]) * (np.linspace(-2,2,50) - (w[0]/_
    ↳w[2])), c='g')
plt.axis('equal')
plt.axis([-1.25, 1.25, -1.25, 1.25])
plt.show()

```

Converged in 1800 steps with w = [0. 0.49997488 -0.50002512]



6 f

Modify the experiment such that $\mathbf{x}_n \in \mathbb{R}^{10}$ instead of \mathbb{R}^2 . Run the algorithm on a randomly generated dataset of size 1000. How many updates does the algorithm take to converge?

```
[12]: # generate a dataset of points in 10D space from -1 to 1
N = 1000
x = (np.random.rand(N,10) * 2) - 1

# make a new target function in  $\mathbb{R}^{10}$ 
# create a "random" target function
w_t = np.random.rand(11)
def target(x_n):
    return 1 if np.dot(w_t,np.insert(x_n, 0, 1)) >= 0 else -1

# create vector of y values (labels) with target function
y = [target(x_n) for x_n in x]

# initialize weights to 0
w = np.zeros(11)
idx = 0

while True:
    # pick a random sample
    n = np.random.randint(0,x.shape[0]-1)
    x_n = np.insert(x[n,:], 0, 1)
    y_n = y[n]

    if y_n > 0 and np.dot(w, x_n) < 0:
        w += x_n

    if y_n < 0 and np.dot(w, x_n) >= 0:
        w -= x_n

    if idx % 5000 == 0:
        if test_convergence(w, x, y):
            break
    idx += 1

print(f"Converged in {idx} steps with w = {w}")
```

```
Converged in 225000 steps with w = [ 3.          24.79573418  1.07691329
29.49325043 17.13808812 23.41487054
25.23573745 25.52347617 30.29276203 14.11818992 33.08671431]
```

7 g

Summarize your conclusions regarding the accuracy and running time of the algorithm as a function of N (the number of data points) and d (the number of dimensions).

As we increase the number of data points, N , we are able to more closely learn the target function. This is shown when we compare the results from the examples with $N=20$ vs $N=100$. In the cases where $N=20$, once we reach convergence, there is often still a significant difference between the true target function and the function we learn. With $N=100$ we get fairly close to the true target function, and at $N=1000$, we learn nearly the same function.

We also note, that as we increase N and d , run time, or the steps required for convergence increases. This can result in a significant slow down if we check for convergence by iterating over all data points at every step. For this reason we opt to only check for convergence every k steps (in the larger models), since a single step is much less costly than checking for convergence.